

Flashcards

8 cards

QUESTION / TERM

What is Retrieval-Augmented Generation (RAG)?

CARD 1 OF 8

ANSWER / DEFINITION

RAG is a collaboration between a writer (LLM) and a library assistant (retriever) that incorporates retrieved excerpts into the final text, making it more accurate, current, and grounded in evidence.

QUESTION / TERM

What are the formal components of RAG?

CARD 2 OF 8

ANSWER / DEFINITION

The formal components of RAG are: retriever, generator, and knowledge base.

QUESTION / TERM

What is the role of the retriever in RAG?

CARD 3 OF 8

ANSWER / DEFINITION

The retriever searches an external knowledge source and returns a small set of relevant chunks.

QUESTION / TERM

What is the role of the generator in RAG?

CARD 4 OF 8

ANSWER / DEFINITION

The generator is a large language model that takes the retrieved chunks as context and produces the final output.

QUESTION / TERM

What is the knowledge base in RAG?

CARD 5 OF 8

ANSWER / DEFINITION

The knowledge base is the corpus of up-to-date information the retriever queries.

QUESTION / TERM

How does RAG differ from 'plain' LLM generation?

CARD 6 OF 8

ANSWER / DEFINITION

In RAG, the LLM augments its reasoning with freshly retrieved evidence, whereas in 'plain' LLM generation, the LLM relies solely on its pre-training.

QUESTION / TERM

What is the step-by-step flow of a classic RAG pipeline?

CARD 7 OF 8

ANSWER / DEFINITION

The step-by-step flow is: 1) user prompt, 2) retriever queries the knowledge base, 3) prompt construction, and 4) generator produces a response.

QUESTION / TERM

What is the purpose of the fusion strategy in RAG?

CARD 8 OF 8

ANSWER / DEFINITION

The fusion strategy determines how retrieved passages are combined with the prompt, such as concatenation, hierarchical attention, or specialized 'RAG-token' layers.

