

Retrieval-Augmented Generation (RAG) Quiz

 5/5 (100%)

QUESTION 1

What is the main limitation of a large language model (LLM) without Retrieval-Augmented Generation (RAG)?

- A** It can access real-time information
- B** It has limited training data
- C** Its memory is limited to what it learned during training ✓
- D** It can only generate short responses

 The LLM's memory is limited to its training data and it cannot access or retrieve up-to-date information.

QUESTION 2

What is the role of the retriever in a RAG pipeline?

- A** To generate the final output
- B** To search an external knowledge source and return relevant chunks ✓
- C** To combine retrieved passages with the prompt
- D** To load the knowledge base

 The retriever searches an external knowledge source and returns a small set of relevant chunks.

QUESTION 3

What is the purpose of the knowledge base in a RAG pipeline?

- A** To store the retrieved chunks
- B** To train the LLM
- C** To provide up-to-date information for the retriever to query ✓
- D** To generate the final output

💡 The knowledge base is the corpus of up-to-date information that the retriever queries.

QUESTION 4

How does the generator use the retrieved chunks in a RAG pipeline?

- A** It ignores the retrieved chunks and generates a response solely from its training data
- B** It uses the retrieved chunks as context to produce the final output ✓
- C** It combines the retrieved chunks with the prompt using a fusion strategy
- D** It returns the retrieved chunks as the final output

💡 The generator takes the retrieved chunks as context and produces the final output.

QUESTION 5

What is the result of using RAG compared to 'plain' LLM generation?

- A** Less accurate and less current
- B** More accurate and more current ✓
- C** Less accurate but more current
- D** More accurate but less current

💡 The result of using RAG is more accurate and more current because it augments the LLM's reasoning with freshly retrieved evidence.

Generated by StudyFlow AI